

weather-prediction

May 31, 2023

```
[633]: import numpy as np
import pandas as pd
```

```
[634]: df=pd.read_csv('daily_weather.csv')
```

```
[635]: df.head(3)
```

```
[635]:
```

	number	air_pressure_9am	air_temp_9am	avg_wind_direction_9am	\
0	0	918.060000	74.822000	271.100000	
1	1	917.347688	71.403843	101.935179	
2	2	923.040000	60.638000	51.000000	

	avg_wind_speed_9am	max_wind_direction_9am	max_wind_speed_9am	\
0	2.080354	295.400000	2.863283	
1	2.443009	140.471548	3.533324	
2	17.067852	63.700000	22.100967	

	rain_accumulation_9am	rain_duration_9am	relative_humidity_9am	\
0	0.0	0.0	42.420000	
1	0.0	0.0	24.328697	
2	0.0	20.0	8.900000	

	relative_humidity_3pm
0	36.160000
1	19.426597
2	14.460000

```
[636]: df.tail()
```

```
[636]:
```

	number	air_pressure_9am	air_temp_9am	avg_wind_direction_9am	\
1090	1090	918.90	63.104	192.9	
1091	1091	918.71	49.568	241.6	
1092	1092	916.60	71.096	189.3	
1093	1093	912.60	58.406	172.7	
1094	1094	921.53	77.702	97.1	

	avg_wind_speed_9am	max_wind_direction_9am	max_wind_speed_9am	\
--	--------------------	------------------------	--------------------	---

1090	3.869906	207.3	5.212070
1091	1.811921	227.4	2.371156
1092	3.064608	200.8	3.892276
1093	3.825167	189.1	4.764682
1094	3.265932	125.9	4.451511

	rain_accumulation_9am	rain_duration_9am	relative_humidity_9am	\
1090	0.0	0.0	26.02	
1091	0.0	0.0	90.35	
1092	0.0	0.0	45.59	
1093	0.0	0.0	64.84	
1094	0.0	0.0	14.56	

	relative_humidity_3pm
1090	38.18
1091	73.34
1092	52.31
1093	58.28
1094	15.10

```
[637]: df.shape
```

```
[637]: (1095, 11)
```

```
[638]: df.isnull().sum()
```

```
[638]: number          0
air_pressure_9am      3
air_temp_9am          5
avg_wind_direction_9am 4
avg_wind_speed_9am    3
max_wind_direction_9am 3
max_wind_speed_9am    4
rain_accumulation_9am  6
rain_duration_9am      3
relative_humidity_9am  0
relative_humidity_3pm  0
dtype: int64
```

```
[639]: df.dropna(inplace=True)
```

```
[640]: df.isnull().sum()
```

```
[640]: number          0
air_pressure_9am      0
air_temp_9am          0
avg_wind_direction_9am 0
```

```

avg_wind_speed_9am      0
max_wind_direction_9am  0
max_wind_speed_9am      0
rain_accumulation_9am   0
rain_duration_9am       0
relative_humidity_9am   0
relative_humidity_3pm    0
dtype: int64

```

```
[641]: df.head(2)
```

```

[641]:   number  air_pressure_9am  air_temp_9am  avg_wind_direction_9am  \
0        0      918.060000    74.822000      271.100000
1        1      917.347688    71.403843      101.935179

      avg_wind_speed_9am  max_wind_direction_9am  max_wind_speed_9am  \
0          2.080354      295.400000      2.863283
1          2.443009      140.471548      3.533324

      rain_accumulation_9am  rain_duration_9am  relative_humidity_9am  \
0                0.0          0.0      42.420000
1                0.0          0.0      24.328697

      relative_humidity_3pm
0          36.160000
1          19.426597

```

```
[642]: df.drop(columns='number',inplace=True)
```

```
[643]: df
```

```

[643]:   air_pressure_9am  air_temp_9am  avg_wind_direction_9am  \
0      918.060000    74.822000      271.100000
1      917.347688    71.403843      101.935179
2      923.040000    60.638000      51.000000
3      920.502751    70.138895      198.832133
4      921.160000    44.294000      277.800000
...          ...          ...          ...
1090     918.900000    63.104000      192.900000
1091     918.710000    49.568000      241.600000
1092     916.600000    71.096000      189.300000
1093     912.600000    58.406000      172.700000
1094     921.530000    77.702000      97.100000

      avg_wind_speed_9am  max_wind_direction_9am  max_wind_speed_9am  \
0          2.080354      295.400000      2.863283
1          2.443009      140.471548      3.533324

```

2	17.067852	63.700000	22.100967
3	4.337363	211.203341	5.190045
4	1.856660	136.500000	2.863283
...
1090	3.869906	207.300000	5.212070
1091	1.811921	227.400000	2.371156
1092	3.064608	200.800000	3.892276
1093	3.825167	189.100000	4.764682
1094	3.265932	125.900000	4.451511

	rain_accumulation_9am	rain_duration_9am	relative_humidity_9am	\
0	0.0	0.0	42.420000	
1	0.0	0.0	24.328697	
2	0.0	20.0	8.900000	
3	0.0	0.0	12.189102	
4	8.9	14730.0	92.410000	
...	
1090	0.0	0.0	26.020000	
1091	0.0	0.0	90.350000	
1092	0.0	0.0	45.590000	
1093	0.0	0.0	64.840000	
1094	0.0	0.0	14.560000	

	relative_humidity_3pm
0	36.160000
1	19.426597
2	14.460000
3	12.742547
4	76.740000
...	...
1090	38.180000
1091	73.340000
1092	52.310000
1093	58.280000
1094	15.100000

[1064 rows x 10 columns]

[644]: `df.info()`

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 1064 entries, 0 to 1094
Data columns (total 10 columns):
#   Column              Non-Null Count  Dtype
---  -
0   air_pressure_9am    1064 non-null   float64
1   air_temp_9am        1064 non-null   float64
```

```

2  avg_wind_direction_9am  1064 non-null  float64
3  avg_wind_speed_9am      1064 non-null  float64
4  max_wind_direction_9am  1064 non-null  float64
5  max_wind_speed_9am      1064 non-null  float64
6  rain_accumulation_9am   1064 non-null  float64
7  rain_duration_9am       1064 non-null  float64
8  relative_humidity_9am   1064 non-null  float64
9  relative_humidity_3pm   1064 non-null  float64
dtypes: float64(10)
memory usage: 91.4 KB

```

```
[645]: df.describe()
```

```

[645]:      air_pressure_9am  air_temp_9am  avg_wind_direction_9am  \
count      1064.000000    1064.000000          1064.000000
mean        918.903180      65.022609          142.306756
std           3.179040      11.168033           69.149472
min          907.990000      36.752000          15.500000
25%          916.595376      57.398000          65.979244
50%          918.942281      65.778479          165.937461
75%          921.169054      73.530872          191.100000
max          929.320000      98.906000          343.400000

      avg_wind_speed_9am  max_wind_direction_9am  max_wind_speed_9am  \
count      1064.000000          1064.000000          1064.000000
mean         5.485793           148.480424           6.999714
std          4.534427           67.154911           5.590790
min          0.693451           28.900000           1.185578
25%          2.245529           76.335351           3.064608
50%          3.869906          176.350000           4.943637
75%          7.264463          201.125000           8.747888
max          23.554978          312.200000           29.840780

      rain_accumulation_9am  rain_duration_9am  relative_humidity_9am  \
count      1064.000000          1064.000000          1064.000000
mean         0.182023          266.393697           34.077440
std          1.534493          1503.092216          25.356668
min          0.000000           0.000000           6.090000
25%          0.000000           0.000000          15.093365
50%          0.000000           0.000000          23.135000
75%          0.000000           0.000000          44.660000
max          24.020000          17704.000000          92.620000

      relative_humidity_3pm
count      1064.000000
mean        35.148381
std         22.365475

```

```

min          5.300000
25%         17.360468
50%         24.371286
75%         51.922500
max          92.250000


```

```
[646]: pie=df['air_pressure_9am'].value_counts().head(10)
```

```
[647]: import matplotlib.pyplot as plt
plt.pie(pie,autopct='%1.1f%%')
plt.title('Top 10 air_pressure_9am in percentage')
```

```
[647]: Text(0.5, 1.0, 'Top 10 air_pressure_9am in percentage')
```

Top 10 air_pressure_9am in percentage



```
[648]: pie
```

```
[648]: 918.60    3
       915.20    2
       917.60    2
       917.01    2
       919.30    2
       925.30    2
```

```

921.34    2
923.60    2
918.71    2
916.10    2
Name: air_pressure_9am, dtype: int64

```

```
[649]: pie2=df['air_pressure_9am'].value_counts().tail(10)
```

```
[650]: plt.pie(pie2,autopct='%1.1f%%')
plt.title('Last 10 air_pressure_ in percentage')
```

```
[650]: Text(0.5, 1.0, 'Last 10 air_pressure_ in percentage')
```

Last 10 air_pressure_ in percentage



```
[651]: df.head()
```

```
[651]:
```


	air_pressure_9am	air_temp_9am	avg_wind_direction_9am	avg_wind_speed_9am	\
0	918.060000	74.822000	271.100000	2.080354	
1	917.347688	71.403843	101.935179	2.443009	
2	923.040000	60.638000	51.000000	17.067852	
3	920.502751	70.138895	198.832133	4.337363	
4	921.160000	44.294000	277.800000	1.856660	

	max_wind_direction_9am	max_wind_speed_9am	rain_accumulation_9am	\
0	295.400000	2.863283	0.0	
1	140.471548	3.533324	0.0	
2	63.700000	22.100967	0.0	
3	211.203341	5.190045	0.0	
4	136.500000	2.863283	8.9	

	rain_duration_9am	relative_humidity_9am	relative_humidity_3pm
0	0.0	42.420000	36.160000
1	0.0	24.328697	19.426597
2	20.0	8.900000	14.460000
3	0.0	12.189102	12.742547
4	14730.0	92.410000	76.740000

```
[652]: import seaborn as sns
plt.figure(figsize=(15,8))
sns.heatmap(df.corr(),annot=True,cmap='rainbow')
```

[652]: <Axes: >



```
[ ]: df.info()
```

```
[653]: x=df.drop(columns=['relative_humidity_3pm'],axis=1)
```



```
[654]: y=df['relative_humidity_3pm']
```

```
[655]: from sklearn.model_selection import train_test_split
```

```
[656]: x_train,x_test,y_train,y_test=train_test_split(x,y,test_size=0.2)
```

```
[657]: df.shape
```

```
[657]: (1064, 10)
```

```
[658]: x_train.shape,x_test.shape,y_train.shape,y_test.shape
```

```
[658]: ((851, 9), (213, 9), (851,), (213,))
```

```
[659]: from sklearn.tree import DecisionTreeRegressor
```

```
[660]: model= DecisionTreeRegressor()
```

```
[661]: model.fit(x_train,y_train)
```

```
[661]: DecisionTreeRegressor()
```

```
[662]: model.predict(x_test)
```

```
[662]: array([[21.73      , 39.52      , 7.30218051, 71.65      , 20.3686575 ,
         44.8       , 71.9       , 32.97      , 25.14      , 44.82      ,
        19.60693827, 13.77      , 77.87      , 21.86553243, 44.27      ,
        19.52606277, 91.16      , 23.8       , 17.33744917, 14.41069359,
        17.75      , 53.02      , 13.48750448, 14.05      , 68.05      ,
        54.77      , 15.25344497, 12.69854598, 18.65026305, 44.85      ,
        21.28846893, 24.37743921, 10.376058   , 14.05374792, 69.67      ,
        44.82      , 54.77      , 56.25      , 48.38      , 14.27588983,
        88.16      , 22.60694278, 8.03       , 13.21      , 53.17      ,
        13.55961808, 13.62      , 12.52      , 75.97      , 71.9       ,
         8.62      , 14.15673274, 14.05      , 13.21      , 20.52      ,
        13.48      , 66.57      , 21.77      , 22.18990678, 84.39      ,
        16.87177039, 23.03597591, 88.67      , 30.57      , 10.376058   ,
        19.59      , 15.52      , 43.96      , 46.6       , 13.77      ,
        42.95      , 62.81      , 84.39      , 12.36298408, 64.52      ,
        14.64990936, 74.1       , 20.9       , 18.27451897, 13.21      ,
        14.15673274, 19.92657856, 51.1       , 53.23      , 28.82      ,
        55.52      , 71.9       , 17.45      , 49.39      , 7.52       ,
        68.39      , 44.59      , 7.94640849, 54.77      , 19.2       ,
        17.71      , 53.84      , 13.52559044, 33.19      , 20.04530549,
        45.19      , 21.03831789, 15.25344497, 22.72      , 19.23      ,
        30.       , 20.21      , 79.38      , 19.52606277, 15.64703999,
        17.3342963 , 18.9414337 , 16.68      , 19.04826571, 79.09      ,
```

```

11.43602921, 38.85      , 17.76      , 18.25795332, 27.45      ,
28.82      , 21.15      , 90.06      , 15.46261112, 76.94      ,
73.65      , 40.91      , 47.58      , 8.76931379, 23.8      ,
44.27      , 19.04826571, 23.8      , 52.12      , 51.52      ,
36.55      , 12.15540564, 21.28846893, 37.8      , 14.22808216,
68.27      , 11.10544045, 91.22      , 15.87360525, 44.82      ,
22.77      , 21.28846893, 74.1      , 63.44      , 21.73      ,
11.24      , 9.98407134, 42.95      , 48.81      , 69.19      ,
20.9      , 19.68      , 19.23      , 28.17      , 48.98      ,
42.22      , 16.81584367, 28.12      , 53.02      , 21.16905026,
13.17      , 19.30666431, 69.74      , 27.45      , 28.43      ,
78.51      , 43.96      , 13.35131297, 28.82      , 22.16642603,
23.11818221, 14.46      , 27.45      , 22.72      , 68.39      ,
56.93      , 68.05      , 17.63      , 51.1      , 14.05      ,
22.4      , 42.22      , 22.78455864, 17.28      , 76.94      ,
24.37743921, 20.71      , 17.76      , 91.22      , 19.2      ,
39.21      , 12.01883617, 53.02      , 7.94640849, 85.96      ,
22.60694278, 26.94      , 48.38      , 27.88      , 59.69      ,
12.52      , 17.02234251, 25.11      , 16.66230871, 80.61      ,
13.09702532, 52.97      , 76.7      ]

```

```
[663]: model.score(x_test,y_test)
```

```
[663]: 0.7416526810297336
```

```
[664]: df.head()
```

```

[664]:   air_pressure_9am  air_temp_9am  avg_wind_direction_9am  avg_wind_speed_9am  \
0      918.060000    74.822000          271.100000          2.080354
1      917.347688    71.403843          101.935179          2.443009
2      923.040000    60.638000           51.000000          17.067852
3      920.502751    70.138895          198.832133           4.337363
4      921.160000    44.294000          277.800000           1.856660

      max_wind_direction_9am  max_wind_speed_9am  rain_accumulation_9am  \
0      295.400000          2.863283          0.0
1      140.471548          3.533324          0.0
2       63.700000          22.100967          0.0
3      211.203341           5.190045          0.0
4      136.500000           2.863283          8.9

      rain_duration_9am  relative_humidity_9am  relative_humidity_3pm
0           0.0          42.420000          36.160000
1           0.0          24.328697          19.426597
2          20.0           8.900000          14.460000
3           0.0          12.189102          12.742547
4         14730.0          92.410000          76.740000

```

```

[665]: x=df.
        ↪drop(columns=['relative_humidity_3pm','max_wind_speed_9am','rain_accumulation_9am','avg_wind_speed_9am'])
        y=df['relative_humidity_3pm']
        x_train,x_test,y_train,y_test=train_test_split(x,y,test_size=0.2)

[699]: mode_r1= DecisionTreeRegressor(max_leaf_nodes=10)

[700]: mode_r1.fit(x_train,y_train)

[700]: DecisionTreeRegressor(max_leaf_nodes=10)

[701]: mode_r1.score(x_test,y_test)

[701]: 0.8142315832148777

[702]: from sklearn.preprocessing import StandardScaler

[703]: scale=StandardScaler()

[704]: scale.fit(x_train,y_train)

[704]: StandardScaler()

[705]: x_test_scale=scale.fit_transform(x_test)
        x_train_scale=scale.fit_transform(x_train)

[706]: mode_r1.fit(x_train_scale,y_train)

[706]: DecisionTreeRegressor(max_leaf_nodes=10)

[707]: mode_r1.score(x_test_scale,y_test)

[707]: 0.7961597898639061

[708]: from sklearn.ensemble import RandomForestRegressor

[709]: model2=RandomForestRegressor()

[710]: model2.fit(x_train,y_train)

[710]: RandomForestRegressor()

[712]: model2.score(x_test_scale,y_test)

[712]: 0.821154255074222

[711]: model2.fit(x_train_scale,y_train)

```

```
[711]: RandomForestRegressor()
```

```
[713]: model2.score(x_test_scale,y_test)
```

```
[713]: 0.821154255074222
```

```
[715]: from sklearn.linear_model import LinearRegression
```

```
[716]: model3=LinearRegression()
```

```
[717]: model3.fit(x_train,y_train)
```

```
[717]: LinearRegression()
```

```
[718]: model3.score(x_test,y_test)
```

```
[718]: 0.8117287686999615
```

```
[719]: model3.fit(x_train_scale,y_train)
```

```
[719]: LinearRegression()
```

```
[720]: model3.score(x_test_scale,y_test)
```

```
[720]: 0.8021842202929085
```

```
[721]: from sklearn.preprocessing import PolynomialFeatures
```

```
[722]: from sklearn.pipeline import make_pipeline
```

```
[723]: model_0 = make_pipeline(PolynomialFeatures(degree=1), LinearRegression())
```

```
[724]: model_0.fit(x_train, y_train)
```

```
[724]: Pipeline(steps=[('polynomialfeatures', PolynomialFeatures(degree=1)),  
                  ('linearregression', LinearRegression())])
```

```
[725]: model_0.score(x_test, y_test)
```

```
[725]: 0.8117287686999659
```

```
[726]: model_01 = make_pipeline(PolynomialFeatures(degree=2), DecisionTreeRegressor())
```

```
[727]: model_01.fit(x_train, y_train)
```

```
[727]: Pipeline(steps=[('polynomialfeatures', PolynomialFeatures()),  
                  ('decisiontreeregressor', DecisionTreeRegressor())])
```

```
[728]: model_01.score(x_test, y_test)
```

```
[728]: 0.7226019380451649
```

```
[760]: from sklearn.pipeline import Pipeline
from sklearn.preprocessing import StandardScaler
from sklearn.feature_selection import SelectKBest
from sklearn.ensemble import RandomForestClassifier

pipe = Pipeline([
    ('scaler', StandardScaler()),
    ('kbest', SelectKBest(k=6)),
    ('rf', RandomForestRegressor(n_estimators=10))
])
```

```
[761]: pipe.fit(x_train, y_train)
```

```
C:\Users\User\AppData\Roaming\Python\Python310\site-
packages\sklearn\feature_selection\_univariate_selection.py:113: RuntimeWarning:
divide by zero encountered in divide
    f = msb / msd
```

```
[761]: Pipeline(steps=[('scaler', StandardScaler()), ('kbest', SelectKBest(k=6)),
                        ('rf', RandomForestRegressor(n_estimators=10))])
```

```
[762]: pipe.score(x_test, y_test)
```

```
[762]: 0.8387840758392073
```

```
[763]: from sklearn.pipeline import Pipeline
from sklearn.preprocessing import StandardScaler
from sklearn.feature_selection import SelectKBest
from sklearn.ensemble import RandomForestClassifier

pipe = Pipeline([
    ('scaler', StandardScaler()),
    ('kbest', SelectKBest(k=6)),
    ('rf', LinearRegression())
])
```

```
[764]: pipe.fit(x_train, y_train)
```

```
C:\Users\User\AppData\Roaming\Python\Python310\site-
packages\sklearn\feature_selection\_univariate_selection.py:113: RuntimeWarning:
divide by zero encountered in divide
    f = msb / msd
```

```
[764]: Pipeline(steps=[('scaler', StandardScaler()), ('kbest', SelectKBest(k=6)),  
                        ('rf', LinearRegression())])
```

```
[765]: pipe.score(x_test, y_test)
```

```
[765]: 0.8117287686999617
```

```
[873]: from sklearn.pipeline import Pipeline  
from sklearn.ensemble import RandomForestClassifier  
from sklearn.model_selection import RandomizedSearchCV  
from sklearn.preprocessing import StandardScaler
```

```
# Define the pipeline  
pipeline = Pipeline([  
    ('scaler', StandardScaler()),  
    ('classifier', RandomForestRegressor()),  
)
```

```
# Define the parameter grid to search over  
param_grid = {  
    'classifier__n_estimators': [100, 200, 300],  
    'classifier__max_depth': [5, 10, 15, 20],  
    'classifier__min_samples_split': [2, 5, 10],  
    'classifier__min_samples_leaf': [1, 2, 4],  
}
```

```
# Define the RandomizedSearchCV object  
search = RandomizedSearchCV(  
    pipeline,  
    param_grid,  
    n_iter=2,  
    cv=2,  
    verbose=5,  
    n_jobs=-1  
)
```

```
[874]: search.fit(x_train, y_train)
```

Fitting 2 folds for each of 2 candidates, totalling 4 fits

```
[874]: RandomizedSearchCV(cv=2,  
                        estimator=Pipeline(steps=[('scaler', StandardScaler()),  
                                                  ('classifier',  
                                                   RandomForestRegressor())]),  
                        n_iter=2, n_jobs=-1,  
                        param_distributions={'classifier__max_depth': [5, 10, 15,  
                                                                      20],
```

```
'classifier__min_samples_leaf': [1, 2,
                                  4],
'classifier__min_samples_split': [2, 5,
                                   10],
'classifier__n_estimators': [100, 200,
                               300]},
```

```
verbose=5)
```


```
[881]: search.score(x_test, y_test)
```

```
[881]: 0.8409377024435153
```

```
[900]: humidity_3pm = df.relative_humidity_3pm[:213]
```

```
[901]: plt.bar(humidity_3pm, bar)
```

```
[901]: <BarContainer object of 213 artists>
```



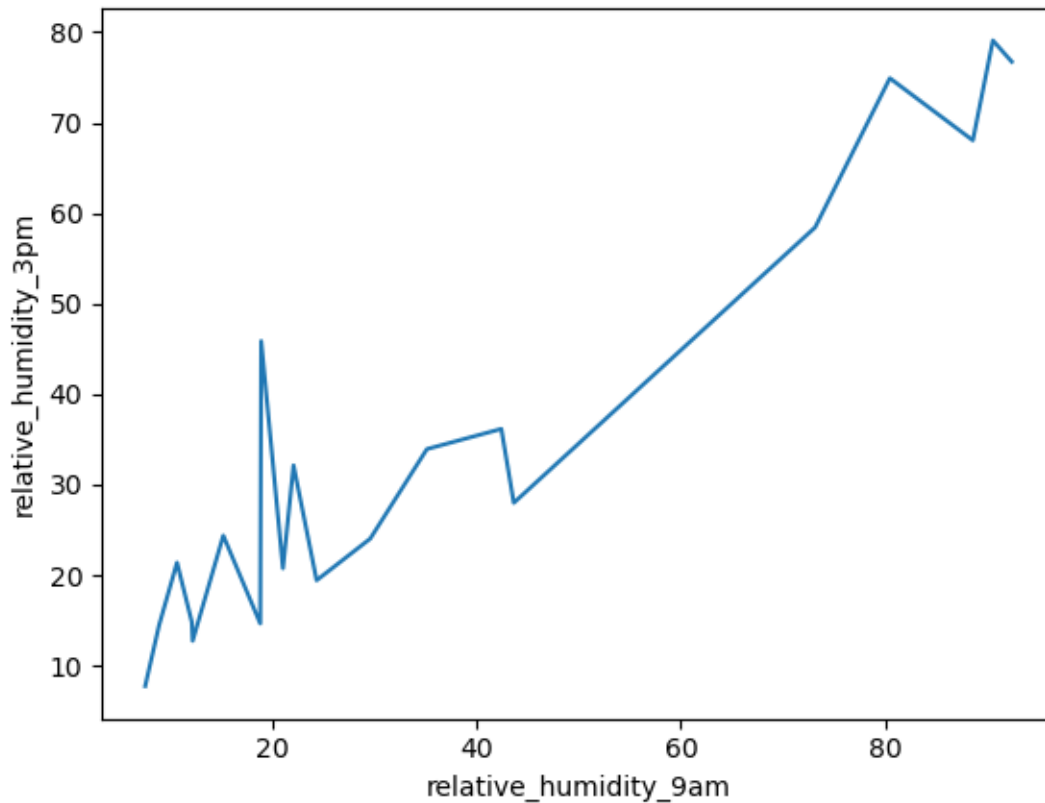
```
[892]: df.head(2)
```

```
[892]:   air_pressure_9am  air_temp_9am  avg_wind_direction_9am  avg_wind_speed_9am  \
0      918.060000    74.822000    271.100000    2.080354
```

1	917.347688	71.403843	101.935179	2.443009
	max_wind_direction_9am	max_wind_speed_9am	rain_accumulation_9am	\
0	295.400000	2.863283	0.0	
1	140.471548	3.533324	0.0	
	rain_duration_9am	relative_humidity_9am	relative_humidity_3pm	
0	0.0	42.420000	36.160000	
1	0.0	24.328697	19.426597	

```
[910]: first=df.relative_humidity_9am.head(20)
second=df.relative_humidity_3pm.head(20)
sns.lineplot(x=first,y=second)
```

```
[910]: <Axes: xlabel='relative_humidity_9am', ylabel='relative_humidity_3pm'>
```



```
[933]: import seaborn as sns
import matplotlib.pyplot as plt

first = df.relative_humidity_9am.head(20)
second = df.relative_humidity_3pm.head(20)
```



```

data = pd.DataFrame({'9am': first, '3pm': second})

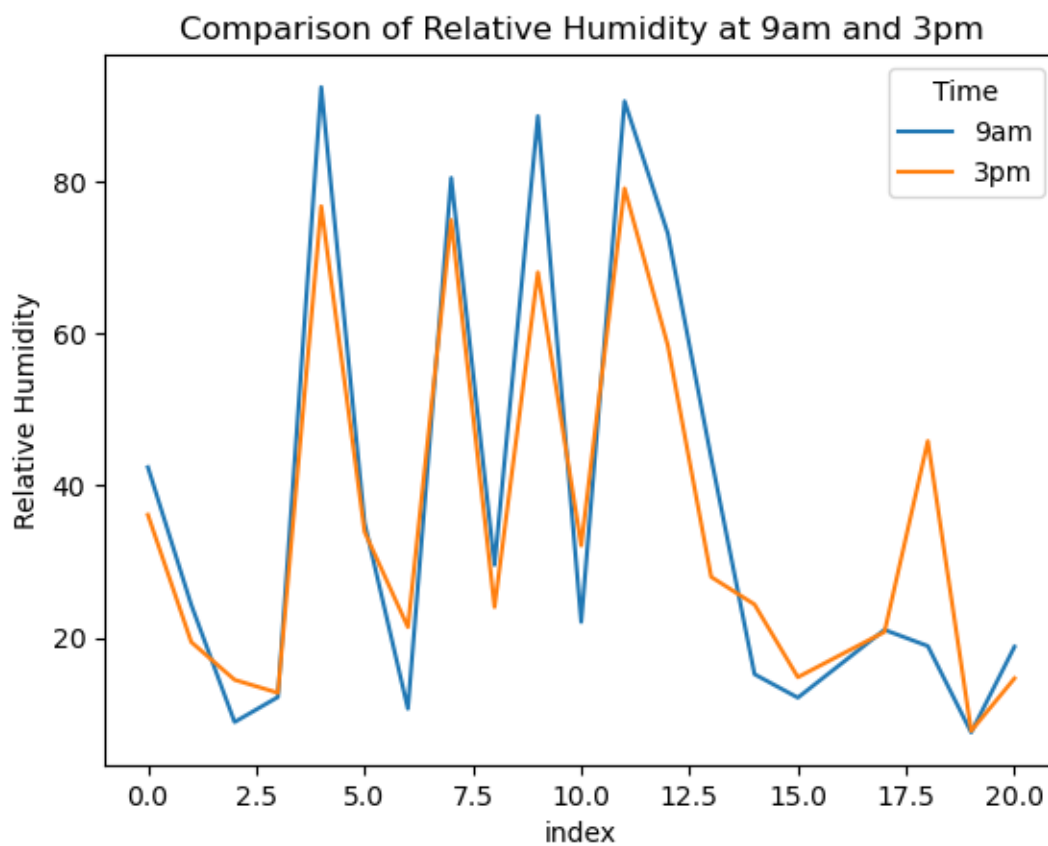
data = data.reset_index()

data = data.melt('index', var_name='Time', value_name='Relative Humidity')

sns.lineplot(data=data, x='index', y='Relative Humidity', hue='Time')

plt.ylabel('Relative Humidity')
plt.title('Comparison of Relative Humidity at 9am and 3pm')
plt.show()

```



```
[927]: columns=['relative_humidity_3pm', 'relative_humidity_9am']
```

```
[928]: df.columns
```

```
[928]:      relative_humidity_3pm  relative_humidity_9am
0          36.160000          42.420000
1          19.426597          24.328697
2          14.460000           8.900000
3          12.742547          12.189102
4          76.740000          92.410000
...
1090         38.180000          26.020000
1091         73.340000          90.350000
1092         52.310000          45.590000
1093         58.280000          64.840000
1094         15.100000          14.560000

[1064 rows x 2 columns]
```

```
[915]: df.head(2)
```

```
[915]:      air_pressure_9am  air_temp_9am  avg_wind_direction_9am  avg_wind_speed_9am \
0          918.060000      74.822000          271.100000          2.080354
1          917.347688      71.403843          101.935179          2.443009

      max_wind_direction_9am  max_wind_speed_9am  rain_accumulation_9am \
0          295.400000          2.863283          0.0
1          140.471548          3.533324          0.0

      rain_duration_9am  relative_humidity_9am  relative_humidity_3pm
0              0.0          42.420000          36.160000
1              0.0          24.328697          19.426597
```

```
[924]: df.columns
```

```
[924]: Index(['air_pressure_9am', 'air_temp_9am', 'avg_wind_direction_9am',
        'avg_wind_speed_9am', 'max_wind_direction_9am', 'max_wind_speed_9am',
        'rain_accumulation_9am', 'rain_duration_9am', 'relative_humidity_9am',
        'relative_humidity_3pm'],
        dtype='object')
```

```
[935]: pred=search.predict(x_test)
```

```
[936]: pred
```

```
[936]: array([17.57243845, 23.77754529, 16.95008649, 16.23809775, 11.91841954,
        16.4803811 , 48.49464826, 26.46887882, 14.47584603,  9.40653724,
        24.0250184 , 16.19651861,  9.519366 , 31.14660244, 69.26240659,
        14.94165528, 58.35702346, 53.68780015, 11.76686343, 14.65498221,
```

12.97874143, 21.34095256, 66.97581498, 17.46205321, 74.70343472,
52.00939102, 16.82342456, 36.76437211, 21.33642087, 16.60116126,
15.23191178, 23.87241921, 15.22239994, 16.57384305, 39.86091483,
16.65948581, 59.2497227 , 16.80865784, 18.5235033 , 63.3059451 ,
10.76640843, 15.02593408, 18.54318033, 66.38477922, 61.49880682,
20.5437302 , 48.68767859, 72.14201337, 44.49042717, 10.70604985,
78.60939765, 62.93346575, 11.48115603, 72.91516903, 62.08392586,
76.45409682, 44.29057488, 20.62282615, 20.42891192, 64.02294837,
16.70774834, 15.17955118, 11.71364036, 17.64697472, 45.83503716,
27.29131362, 18.4208507 , 19.48561187, 44.1288649 , 52.78051355,
29.30849223, 18.11050788, 51.59674448, 44.21592687, 48.7511616 ,
79.9215127 , 47.32151614, 53.4624155 , 59.25726656, 26.3796573 ,
18.75162672, 67.76444309, 17.8541254 , 26.61438934, 59.79107825,
19.18905097, 23.18612873, 57.61944194, 29.04241619, 23.67777368,
49.81289783, 16.66850841, 49.45989075, 14.50549918, 66.86044206,
26.67416756, 67.17034298, 15.8326133 , 23.36276116, 26.57944828,
70.94165738, 17.45527987, 80.40403665, 32.24090993, 17.9286429 ,
38.19385379, 24.75122554, 15.26457471, 66.25852718, 25.6905802 ,
12.76291363, 21.13518022, 24.16921558, 12.96065411, 53.60239179,
36.54354668, 63.18385587, 45.50531681, 77.13383103, 12.66684859,
22.37806395, 40.47163421, 66.57457705, 61.24944428, 19.27881255,
42.02096509, 17.36823668, 19.05929126, 21.91498768, 31.99101065,
19.39270086, 23.93809152, 30.13030535, 28.64643014, 9.96692614,
49.23248637, 49.35538668, 30.40631521, 16.87972701, 23.68795133,
30.34141264, 17.06456603, 18.88063582, 78.558013 , 47.44369879,
17.03475024, 13.28221761, 20.42708138, 17.26652423, 15.02218844,
11.73889446, 29.1763011 , 26.18290365, 21.51521139, 24.77429214,
50.99727859, 38.44044377, 49.36757333, 17.74998581, 50.54576354,
16.23919355, 33.43181118, 24.81092724, 20.45618373, 25.03577347,
28.68566477, 15.02218844, 12.2432319 , 71.08293024, 21.25306833,
16.17957236, 18.61009884, 16.6559143 , 19.58078439, 77.90475452,
79.53238779, 25.02273207, 43.11147672, 56.14268016, 16.56988981,
34.77261083, 17.44461617, 76.19047319, 43.50783339, 53.83071621,
72.28141966, 54.43222674, 19.19783723, 44.34388574, 18.97128044,
28.72472229, 19.85879655, 21.0861607 , 15.21248029, 78.42215647,
16.62072307, 23.74517324, 24.97602047, 19.70828239, 18.3825252 ,
31.73416282, 9.21498814, 24.20264126, 15.48349726, 60.06214547,
14.91209476, 21.13816854, 71.12341299, 78.77885676, 38.12668596,
20.31708064, 68.62622195, 23.93336637])

[938]: `pred.shape`

[938]: (213,)

[940]: `y_test.shape`

[940]: (213,)

```
[949]: df.describe()
```

```
[949]:
```

	air_pressure_9am	air_temp_9am	avg_wind_direction_9am	\
count	1064.000000	1064.000000	1064.000000	
mean	918.903180	65.022609	142.306756	
std	3.179040	11.168033	69.149472	
min	907.990000	36.752000	15.500000	
25%	916.595376	57.398000	65.979244	
50%	918.942281	65.778479	165.937461	
75%	921.169054	73.530872	191.100000	
max	929.320000	98.906000	343.400000	

	avg_wind_speed_9am	max_wind_direction_9am	max_wind_speed_9am	\
count	1064.000000	1064.000000	1064.000000	
mean	5.485793	148.480424	6.999714	
std	4.534427	67.154911	5.590790	
min	0.693451	28.900000	1.185578	
25%	2.245529	76.335351	3.064608	
50%	3.869906	176.350000	4.943637	
75%	7.264463	201.125000	8.747888	
max	23.554978	312.200000	29.840780	

	rain_accumulation_9am	rain_duration_9am	relative_humidity_9am	\
count	1064.000000	1064.000000	1064.000000	
mean	0.182023	266.393697	34.077440	
std	1.534493	1503.092216	25.356668	
min	0.000000	0.000000	6.090000	
25%	0.000000	0.000000	15.093365	
50%	0.000000	0.000000	23.135000	
75%	0.000000	0.000000	44.660000	
max	24.020000	17704.000000	92.620000	

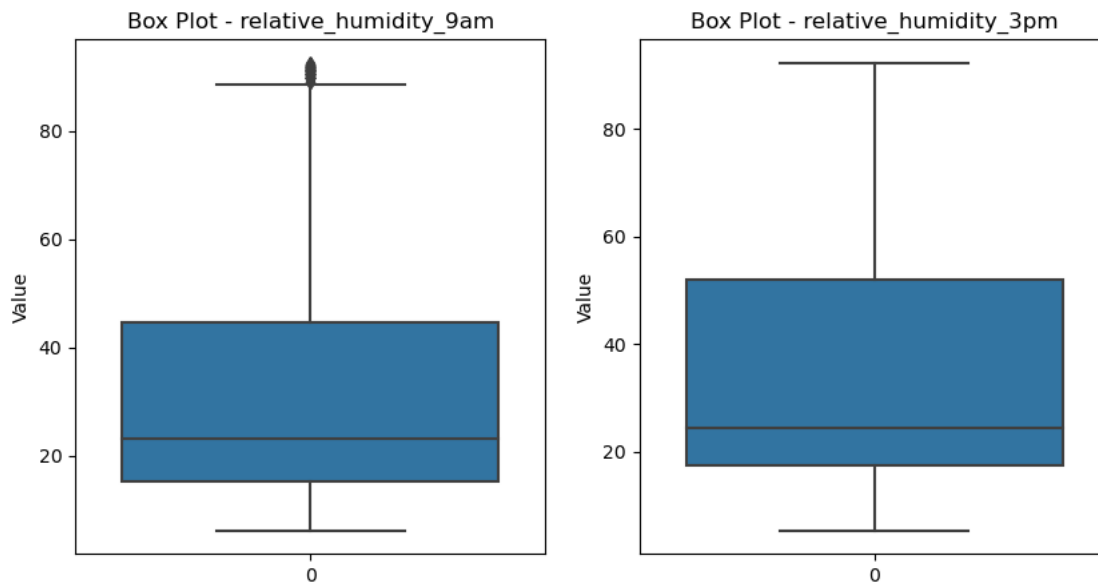
	relative_humidity_3pm
count	1064.000000
mean	35.148381
std	22.365475
min	5.300000
25%	17.360468
50%	24.371286
75%	51.922500
max	92.250000

```
[962]: import seaborn as sns
import matplotlib.pyplot as plt

columns = ['relative_humidity_9am', 'relative_humidity_3pm']
```

```
fig, axes = plt.subplots(nrows=1, ncols=len(columns), figsize=(10, 5))

for i, column in enumerate(columns):
    sns.boxplot(data=df[column], ax=axes[i])
    axes[i].set_ylabel('Value')
    axes[i].set_title(f'Box Plot - {column}')
```



```
[964]: df.head(2)
```

```
[964]:   air_pressure_9am  air_temp_9am  avg_wind_direction_9am  avg_wind_speed_9am \
0         918.060000      74.822000             271.100000          2.080354
1         917.347688      71.403843             101.935179          2.443009

      max_wind_direction_9am  max_wind_speed_9am  rain_accumulation_9am \
0                295.400000          2.863283              0.0
1                140.471548          3.533324              0.0

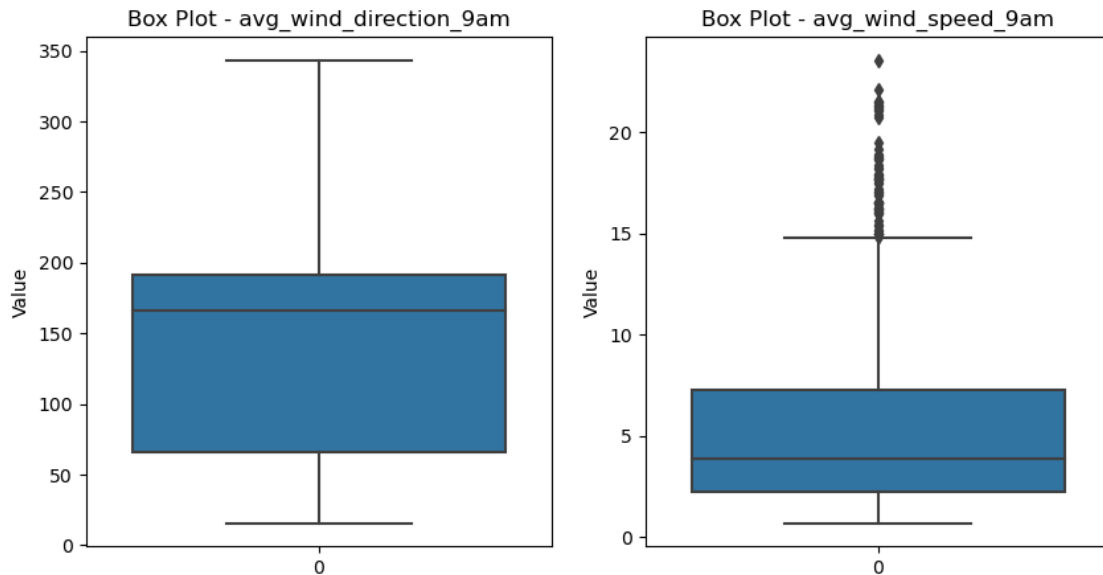
      rain_duration_9am  relative_humidity_9am  relative_humidity_3pm
0                0.0          42.420000          36.160000
1                0.0          24.328697          19.426597
```

```
[965]: import seaborn as sns
import matplotlib.pyplot as plt

columns = ['avg_wind_direction_9am', 'avg_wind_speed_9am']
```

```
fig, axes = plt.subplots(nrows=1, ncols=len(columns), figsize=(10, 5))

for i, column in enumerate(columns):
    sns.boxplot(data=df[column], ax=axes[i])
    axes[i].set_ylabel('Value')
    axes[i].set_title(f'Box Plot - {column}')
```



```
[966]: import pandas as pd
import numpy as np

# Assuming you have a DataFrame called 'df' with multiple columns

# Define a function to detect outliers using Tukey's fences method
def detect_outliers(data):
    Q1 = np.percentile(data, 25)
    Q3 = np.percentile(data, 75)
    IQR = Q3 - Q1
    lower_bound = Q1 - 1.5 * IQR
    upper_bound = Q3 + 1.5 * IQR
    outliers = data[(data < lower_bound) | (data > upper_bound)]
    return outliers

# Iterate over each column and find outliers
for column in df.columns:
```

```

outliers = detect_outliers(df[column])

if outliers.empty:
    print(f"No outliers found in column: {column}")
else:
    print(f"Outliers found in column: {column}")
    print(outliers)
print()

```

Outliers found in column: air_pressure_9am

204 907.99

281 929.32

342 908.42

474 908.97

708 909.19

Name: air_pressure_9am, dtype: float64

Outliers found in column: air_temp_9am

720 98.906

Name: air_temp_9am, dtype: float64

No outliers found in column: avg_wind_direction_9am

Outliers found in column: avg_wind_speed_9am

2 17.067852

61 16.486248

67 17.527932

97 23.554978

104 16.126429

...

1012 18.698476

1017 14.987498

1037 18.653986

1044 17.188835

1078 15.162800

Name: avg_wind_speed_9am, Length: 67, dtype: float64

No outliers found in column: max_wind_direction_9am

Outliers found in column: max_wind_speed_9am

2 22.100967

61 18.029736

67 20.971183

71 19.214314

97 29.840780

...

1012 23.377962

```
1017    19.617964
1037    23.268404
1044    22.195713
1078    19.803578
```

Name: max_wind_speed_9am, Length: 86, dtype: float64

Outliers found in column: rain_accumulation_9am

```
4        8.900
5         0.020
11        0.550
42        1.530
46        0.021
```

...

```
1009    0.010
1011    0.010
1018    0.010
1027    5.272
1059    0.080
```

Name: rain_accumulation_9am, Length: 121, dtype: float64

Outliers found in column: rain_duration_9am

```
2         20.000000
4       14730.000000
5         170.000000
11        1770.000000
27        220.000000
```

...

```
1037     11.024881
1039      6.029673
1044     19.348384
1059    520.000000
1076     10.000000
```

Name: rain_duration_9am, Length: 176, dtype: float64

Outliers found in column: relative_humidity_9am

```
4        92.41
11       90.56
26       92.10
36       89.77
38       91.22
```

...

```
1082     89.80
1083     91.03
1085     92.30
1086     91.11
1091     90.35
```

Name: relative_humidity_9am, Length: 74, dtype: float64

No outliers found in column: relative_humidity_3pm

```
[967]: df.head(2)
```

```
[967]:   air_pressure_9am  air_temp_9am  avg_wind_direction_9am  avg_wind_speed_9am \
0         918.060000      74.822000             271.100000             2.080354
1         917.347688      71.403843             101.935179             2.443009

      max_wind_direction_9am  max_wind_speed_9am  rain_accumulation_9am \
0             295.400000             2.863283             0.0
1             140.471548             3.533324             0.0

      rain_duration_9am  relative_humidity_9am  relative_humidity_3pm
0              0.0             42.420000             36.160000
1              0.0             24.328697             19.426597
```

```
[971]: df.columns
```

```
[971]: Index(['air_pressure_9am', 'air_temp_9am', 'avg_wind_direction_9am',
      'avg_wind_speed_9am', 'max_wind_direction_9am', 'max_wind_speed_9am',
      'rain_accumulation_9am', 'rain_duration_9am', 'relative_humidity_9am',
      'relative_humidity_3pm'],
      dtype='object')
```

```
[979]: df.describe()
```

```
[979]:   air_pressure_9am  air_temp_9am  avg_wind_direction_9am \
count      1064.000000      1064.000000      1064.000000
mean        918.903180       65.022609      142.306756
std          3.179040       11.168033       69.149472
min         907.990000       36.752000       15.500000
25%         916.595376       57.398000       65.979244
50%         918.942281       65.778479      165.937461
75%         921.169054       73.530872      191.100000
max         929.320000       98.906000      343.400000

      avg_wind_speed_9am  max_wind_direction_9am  max_wind_speed_9am \
count      1064.000000      1064.000000      1064.000000
mean         5.485793       148.480424         6.999714
std          4.534427        67.154911         5.590790
min          0.693451       28.900000         1.185578
25%          2.245529       76.335351         3.064608
50%          3.869906      176.350000         4.943637
75%          7.264463      201.125000         8.747888
max         23.554978      312.200000        29.840780
```

	rain_accumulation_9am	rain_duration_9am	relative_humidity_9am \
count	1064.000000	1064.000000	1064.000000
mean	0.182023	266.393697	34.077440
std	1.534493	1503.092216	25.356668
min	0.000000	0.000000	6.090000
25%	0.000000	0.000000	15.093365
50%	0.000000	0.000000	23.135000
75%	0.000000	0.000000	44.660000
max	24.020000	17704.000000	92.620000

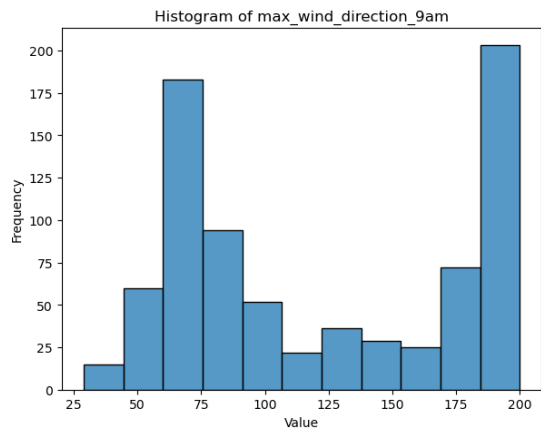
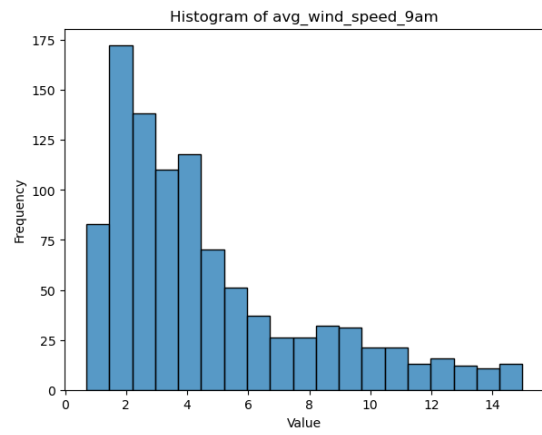

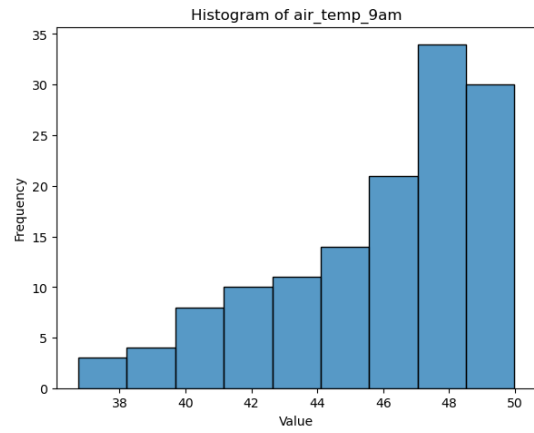
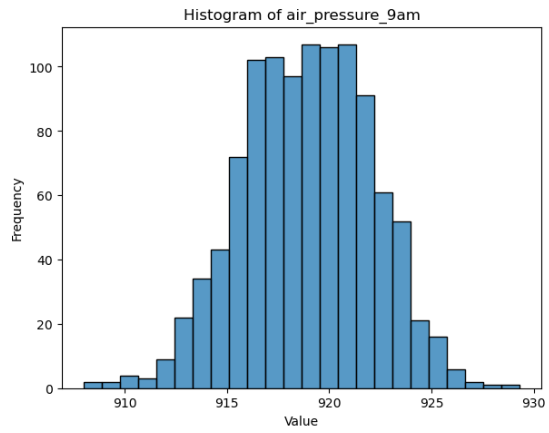
	relative_humidity_3pm
count	1064.000000
mean	35.148381
std	22.365475
min	5.300000
25%	17.360468
50%	24.371286
75%	51.922500
max	92.250000

```
[1000]: import seaborn as sns
import matplotlib.pyplot as plt

air_pressure_9am = df[df['air_pressure_9am'] > 907]['air_pressure_9am']
air_temp_9am = df[df['air_temp_9am'] < 50]['air_temp_9am']
avg_wind_direction_9am = df[df['avg_wind_direction_9am'] <
    ↪200]['avg_wind_direction_9am']
avg_wind_speed_9am = df[df['avg_wind_speed_9am'] < 15]['avg_wind_speed_9am']
max_wind_direction_9am = df[df['max_wind_direction_9am'] <
    ↪200]['max_wind_direction_9am']

show = [air_pressure_9am, air_temp_9am, avg_wind_direction_9am,
    ↪avg_wind_speed_9am, max_wind_direction_9am]
variables = ['air_pressure_9am', 'air_temp_9am', 'avg_wind_direction_9am',
    ↪'avg_wind_speed_9am', 'max_wind_direction_9am']

plt.figure(figsize=(15, 30))
for i, col in enumerate(show):
    plt.subplot(5, 2, i + 1)
    sns.histplot(x=col)
    plt.xlabel('Value')
    plt.ylabel('Frequency')
    plt.title(f'Histogram of {variables[i]}')
```



[]:

[]: